Redstone Catalyst

CONFIDENTIAL

AI INTEGRATION PROPOSAL

Fansbook AI Integration

A comprehensive suite of 6 AI-powered features to automate moderation, verify identities, prevent fraud, and supercharge creator engagement — delivered in 4–5 weeks.

PROJECT	DATE
Fansbook	**March 2026**
CLIENT	PHASE
Fansbook	**AI Integration**

v1.0

01

Executive Summary

Fansbook is requesting the integration of intelligent AI capabilities across its creator-fan platform. Rather than building and training AI models from scratch — which would cost hundreds of thousands of dollars and take years — we will integrate best-in-class, production-ready AI services that deliver identical results at a fraction of the cost. All 6 features will be delivered within 4–5 weeks for a total investment of $8,000.

6	4–5	$8,000	~$50–150
AI FEATURES	WKS / DELIVERY	TOTAL COST	MONTHLY RUNNING COST (CLIENT)

All 6 AI Features

- **ID & Age Verification** — AI-powered identity checks on every new user
- **Content Moderation** — Automatic NSFW/illegal content detection on all uploads
- **Fraud & Chargeback Prevention** — Real-time transaction risk scoring
- **Creator Chat & Engagement AI** — Smart reply suggestions for creator–fan messages
- **Creator Upsell Advisor** — AI-driven insights to help creators grow revenue
- **Support Chat AI** — 24/7 automated customer support with human escalation

Why Not Train Our Own AI?

Training a proprietary AI model from scratch requires millions of dollars in compute, massive labeled datasets, and a team of ML engineers working for 12–24 months. We have the capability to do this, and it remains an option for future phases. However, for this phase, integrating established AI services (AWS, Google Cloud, Anthropic Claude, Stripe) delivers the same — and often superior — results immediately, with far lower risk and cost. The client pays only for what they use.

02

Identity & Age Verification

Every user on Fansbook must be 18+. Currently this relies on a self-declaration popup. This feature replaces that with document-based AI verification — ensuring legal compliance and protecting the platform.

How It Works

1. User uploads a government-issued ID (passport, driving licence, national ID)
2. Onfido / Veriff AI reads and validates the document in under 3 seconds
3. Facial liveness check confirms the person holding the ID is real and present
4. Result (Approved / Rejected / Manual Review) is sent to Fansbook backend
5. Admin dashboard shows verification status per user; manual review queue for edge cases

What Gets Built

- ✓ **ID Upload Flow** — Secure upload screen in the user onboarding journey
- ✓ **Liveness Detection** — Selfie + ID face match with anti-spoofing
- ✓ **Backend Webhook** — Receive and store verification results in the database
- ✓ **Admin Review Panel** — View, approve, or reject flagged verifications
- ✓ **Status Badges** — Verified / Unverified indicators across user profiles
- ✓ **Retry Logic** — Users can re-submit if initial attempt fails

Technology

Onfido API | **Veriff (alternative)** | **Node.js Webhook Handler** | **Prisma DB** | **React Upload UI**

Legal Importance

Without proper age verification, the platform is exposed to significant legal liability in the EU, UK, and US. This feature is the highest priority of the entire integration.

ITEM	DETAIL
DELIVERY TIME	1.5 weeks
DEVELOPMENT COST	$1,200
CLIENT MONTHLY COST	~$1.50 per new verified user

03

AI Content Moderation

Every image and video uploaded to Fansbook is automatically scanned before it goes live. This protects the platform from illegal content, CSAM, extreme violence, and policy violations — without requiring a full-time moderation team.

How It Works

1. Creator uploads image or video
2. Before storing, our backend sends the file to AWS Rekognition / Google Vision AI
3. AI returns a confidence score across categories: explicit, suggestive, violence, etc.
4. Content above threshold is auto-blocked; borderline content is flagged for human review
5. Admin sees a moderation queue with AI scores, can approve or remove flagged content
6. Creator is notified if their content is removed with a reason

What Gets Built

- ✓ **Upload Intercept** — Scan before content is stored or shown publicly
- ✓ **Auto-Block Rules** — Configurable thresholds per content category
- ✓ **Admin Moderation Queue** — Review flagged content with AI confidence scores
- ✓ **Creator Notifications** — Automated alerts when content is removed
- ✓ **Audit Log** — Full history of moderation decisions for compliance
- ✓ **Appeal System** — Creators can submit an appeal for removed content

Technology

AWS Rekognition | **Google Cloud Vision AI (backup)** | **S3 Pre-processing Hook** | **React Admin UI**

ITEM	DETAIL
DELIVERY TIME	1 week
DEVELOPMENT COST	$900
CLIENT MONTHLY COST	~$1 per 1,000 images (effectively near-zero)

04

Fraud & Chargeback Prevention

Chargebacks and payment fraud are among the biggest revenue leaks on creator platforms. This feature uses AI to score every transaction in real time and block suspicious activity before it becomes a dispute.

How It Works

1. Every payment on the platform passes through Stripe Radar (built into Stripe)
2. Radar's ML model — trained on billions of global transactions — scores the risk
3. High-risk transactions are blocked or sent to 3D Secure challenge
4. Custom rules are added on top: velocity checks, geo-anomalies, new account spending limits
5. Admin dashboard shows fraud metrics, blocked payments, and dispute rates
6. Optional: Sift integration for deeper behavioural fraud signals (account takeover, promotion abuse)

What Gets Built

- **Stripe Radar Activation** — Enable and configure built-in fraud ML
- **Custom Rule Set** — Platform-specific rules (new accounts, high-value purchases, VPN detection)
- **Admin Fraud Dashboard** — Blocked transactions, risk scores, dispute tracking
- **Velocity Limits** — Rate-limit suspicious spending patterns
- **Dispute Alerts** — Real-time notifications when chargebacks are filed
- **Reporting** — Monthly fraud summary for the client

Technology

Stripe Radar | Stripe Webhooks | Custom Rule Engine | Node.js | React Admin Dashboard

Fastest ROI of All 6 Features

Stripe Radar is already partially available in the existing Stripe integration. Most of this work is configuration and custom rules. The prevention of even a handful of chargebacks per month will recover this feature's cost entirely.

ITEM	DETAIL
DELIVERY TIME	3–4 days
DEVELOPMENT COST	$700
CLIENT MONTHLY COST	Free (Stripe Radar included) / $0.05 per transaction for premium

05

Creator Chat & Engagement AI

Creators on Fansbook often manage hundreds of fan conversations simultaneously. This feature gives creators an AI assistant inside their message inbox — suggesting smart replies, drafting messages in their tone, and helping them stay engaged with fans without burning out.

How It Works

1. Creator opens a fan conversation in their inbox
2. An AI suggestion bar appears below the message input with 2–3 smart reply options
3. Creator can tap a suggestion to use it as-is, or edit before sending
4. Over time, the AI learns the creator's tone and style (based on their message history)
5. Creators can also type a rough idea and ask AI to polish it into a full message
6. All AI usage is optional — creators can ignore suggestions entirely

What Gets Built

- ✓ **Smart Reply Bar** — 2–3 AI-generated reply suggestions in message composer
- ✓ **Tone Learning** — AI adapts to the creator's writing style over time
- ✓ **Polish Mode** — Type rough idea, AI rewrites as a full engaging message
- ✓ **Context Awareness** — AI reads conversation history for relevant suggestions
- ✓ **Usage Controls** — Creators can enable/disable AI assistance per conversation
- ✓ **Token Usage Tracking** — Admin can monitor AI usage and costs per creator

Technology

Claude API (Anthropic) | **OpenAI GPT-4o (alternative)** | **React Message UI** | **Node.js Middleware** | **Redis Context Cache**

ITEM	DETAIL
DELIVERY TIME	2.5 weeks
DEVELOPMENT COST	$1,800
CLIENT MONTHLY COST	~$20–60/month depending on platform usage

06

Creator Upsell Advisor

Most creators leave money on the table without realising it. This feature deploys an AI advisor inside the creator dashboard that analyses performance data and delivers personalised, actionable recommendations to help creators grow their revenue.

How It Works

1. The system collects creator data: subscriber count, tier pricing, post frequency, engagement rate, earnings history
2. This data is sent to Claude API with a structured prompt built for creator monetisation advice
3. AI generates 3–5 prioritised recommendations specific to that creator
4. Insights are displayed as a card on the creator dashboard, refreshed weekly
5. Examples: “Your top 12 fans haven’t had a new post in 8 days — post today to reduce churn”, “Your $9.99 tier has 80 subscribers — consider a premium $24.99 tier”
6. Admin can view aggregate AI insights across all creators

What Gets Built

- ✓ **Insights Dashboard Card** — Weekly AI recommendations on creator home screen
- ✓ **Data Collection Pipeline** — Aggregate creator stats for AI analysis
- ✓ **Prompt Engineering** — Carefully designed prompts tuned for monetisation advice
- ✓ **Weekly Refresh** — Automated background job to regenerate insights
- ✓ **Admin Overview** — Platform-wide insight trends and creator performance
- ✓ **Push Notifications** — Alert creators when new insights are available

Technology

Claude API | **BullMQ (existing job queue)** | **Node.js** | **React Dashboard Card** | **Prisma Analytics Queries**

ITEM	DETAIL
DELIVERY TIME	2 weeks
DEVELOPMENT COST	$1,600
CLIENT MONTHLY COST	Included in shared Claude API usage (~$20–60/month total with Feature 4)

07

AI Support Chatbot

Support requests consume significant time and cost. This feature deploys a 24/7 AI chatbot that handles the majority of common questions instantly — with seamless escalation to a human agent when needed.

How It Works

1. A chat widget appears on all pages for logged-in users (bottom-right corner)
2. User types a question; the AI searches a knowledge base of FAQs, policies, and platform guides
3. Claude AI generates a helpful, accurate response in plain language
4. If the AI cannot confidently answer, it escalates to a human support agent with full conversation history
5. Admin can add/edit knowledge base articles to keep the AI up to date
6. All conversations are logged; admin can review and improve AI responses over time

What Gets Built

- **Chat Widget UI** — Floating support chat accessible from all pages
- **Knowledge Base** — Structured FAQ and policy articles fed to the AI
- **Claude Integration** — AI response generation with context from knowledge base
- **Human Escalation** — Route to human agent with full chat history when AI is unsure
- **Admin Knowledge Editor** — Add, edit, and manage AI knowledge base articles
- **Conversation Logs** — Full history of all support interactions for QA

Technology

Claude API | React Chat Widget | Node.js WebSocket | Knowledge Base CMS | Admin Panel

Support Cost Reduction

Based on industry data, AI chatbots resolve 60–80% of support tickets without human involvement. For a growing platform like Fansbook, this alone can save the equivalent of 1–2 part-time support staff.

ITEM	DETAIL
DELIVERY TIME	1.5 weeks
DEVELOPMENT COST	$1,800
CLIENT MONTHLY COST	~$20–40/month (Claude API usage)

08

Investment & Timeline

All six features are developed in parallel by our team over 4–5 weeks. The total one-time development investment is $8,000. Ongoing platform running costs are minimal and usage-based.

Development Cost Breakdown

#	FEATURE	DELIVERY	COST
1	ID & Age Verification	Week 1–2	$1,200
2	Content Moderation	Week 1	$900
3	Fraud & Chargeback Prevention	Week 1	$700
4	Creator Chat & Engagement AI	Week 2–4	$1,800
5	Creator Upsell Advisor	Week 3–4	$1,600
6	AI Support Chatbot	Week 3–5	$1,800
	Total	**4–5 Weeks**	**$8,000**

Client Monthly Running Costs (Post-Launch)

SERVICE	PURPOSE	ESTIMATED MONTHLY COST
Onfido / Veriff	ID verification	~$1.50 per new verified user
AWS Rekognition	Content moderation	~$1 per 1,000 images
Stripe Radar	Fraud prevention	Free / $0.05 per transaction
Claude API	Chat AI + Upsell + Support	~$50–100/month
Total	**All AI services**	**~$50–150/month**

4–5 Week Delivery Timeline

WEEK	WORK
Week 1	ID Verification integration, Content Moderation pipeline, Stripe Radar + fraud rules
Week 2	ID Verification admin panel, Creator Chat AI foundation, Upsell Advisor data pipeline
Week 3	Creator Chat AI smart replies + tone learning, Upsell Advisor dashboard card, Support chatbot foundation
Week 4	Support chatbot knowledge base + escalation, Upsell weekly job, polish + QA on all features
Week 5	Final testing, bug fixes, deployment, client handover & documentation

Payment Structure

$2,000 upfront to begin work. Remaining $6,000 paid progressively as features are delivered — you pay for each feature as it is built and handed over, not all at once.

10

Next Steps

1. Client reviews and approves this proposal
2. Development begins — weekly progress updates provided
3. All 6 features delivered, tested, and deployed within 4–5 weeks
4. Handover documentation, admin guides, and 30-day post-launch support included

What's Included in $8,000

Full development of all 6 AI features. Integration with your existing Fansbook codebase. Admin panels for each feature. Deployment to your production server. 30 days of post-launch bug fixes at no additional charge.

COMPANY	PROJECT	TOTAL INVESTMENT
Redstone Catalyst	**Fansbook AI Integration**	**$8,000**